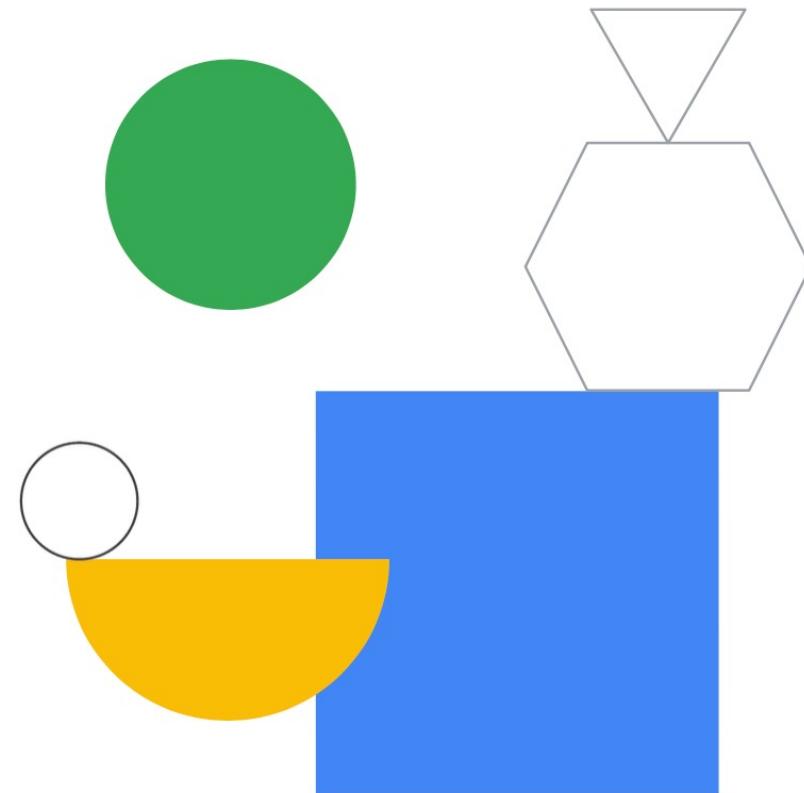


Serverless Messaging with Pub/Sub



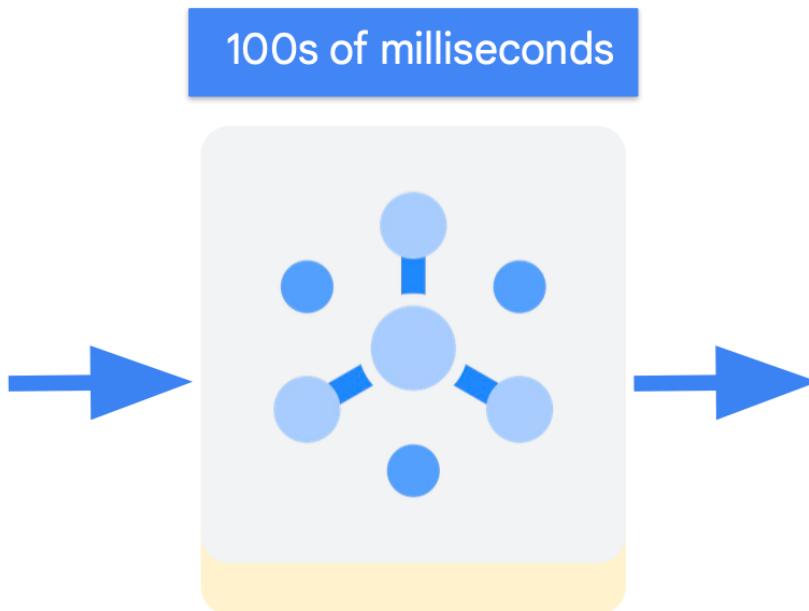
Module agenda

- 01** Introduction to Pub/Sub
 - 02** Pub/Sub Push Versus Pull
 - 03** Publishing with Pub/Sub Code
- 



Introduction to Pub/Sub

Pub/Sub



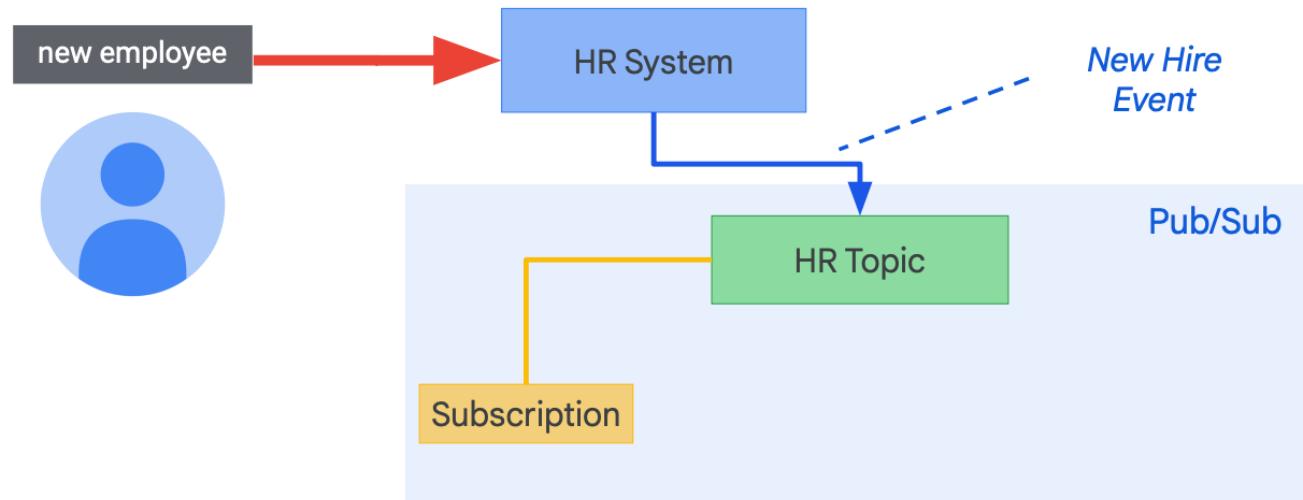
Qualities that Pub/Sub contribute to data engineering solutions:

- ✓ Availability
- ✓ Durability
- ✓ Scalability

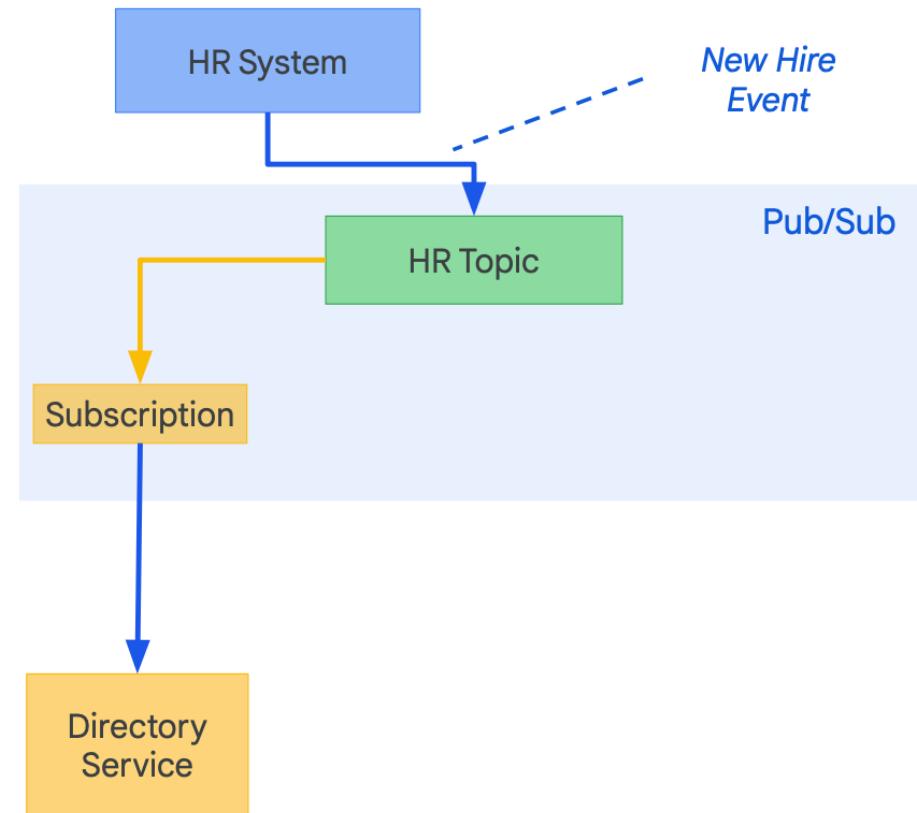
Example of a Pub/Sub application



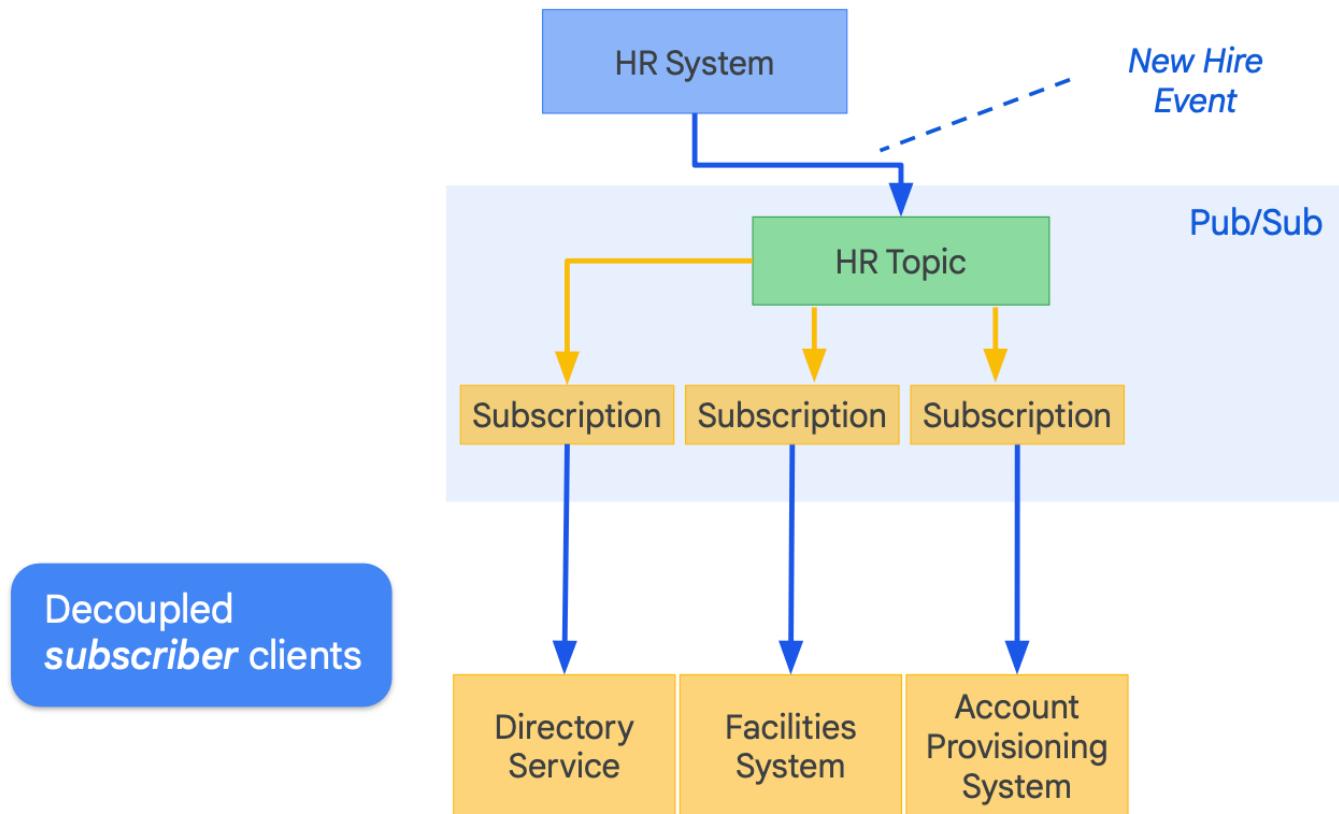
A new employee arrives causing a new hire event



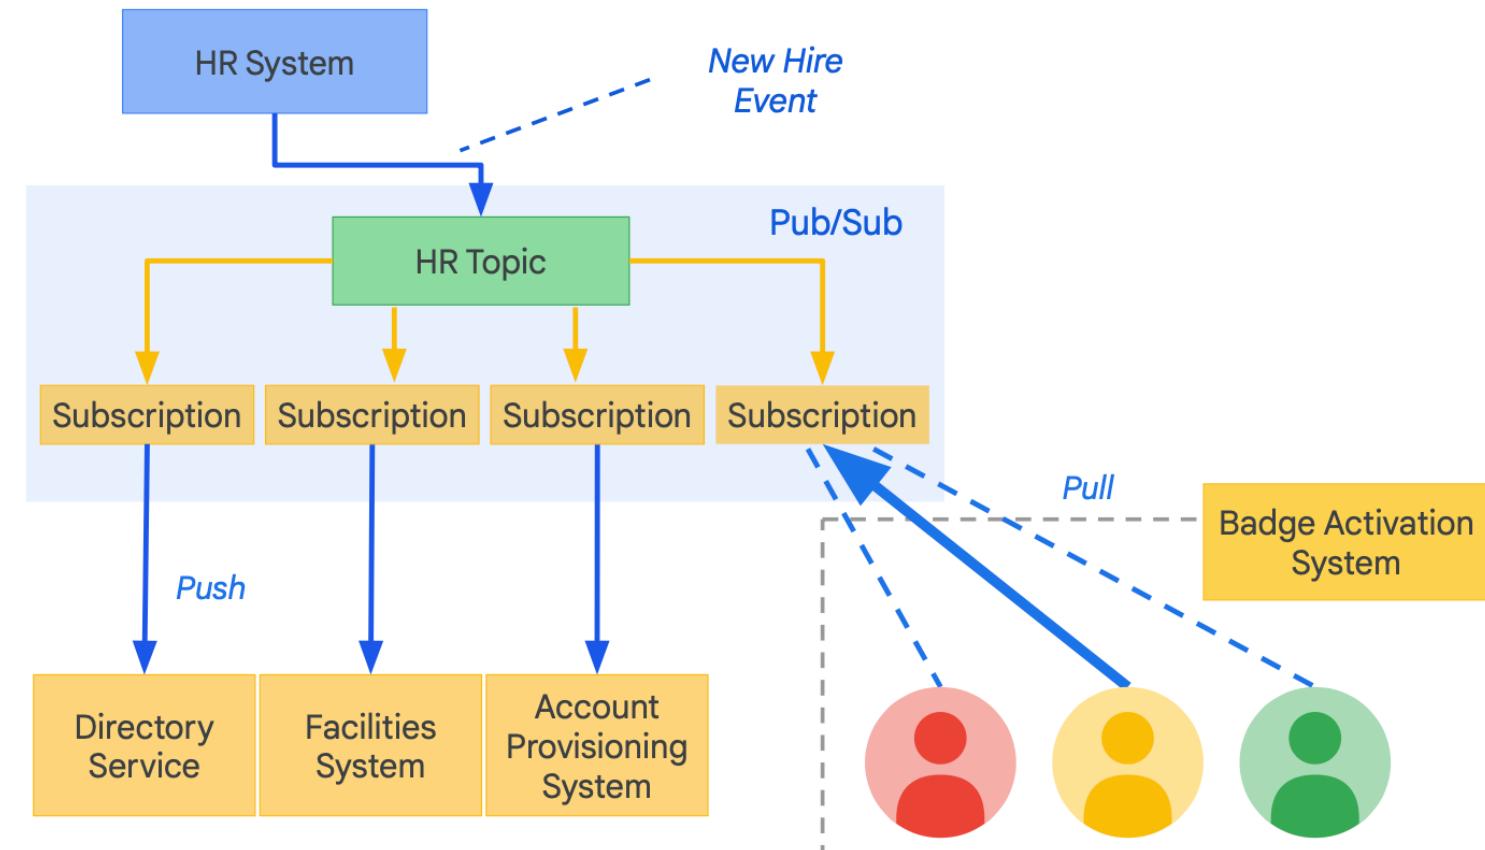
The message is sent from Topic to Subscription



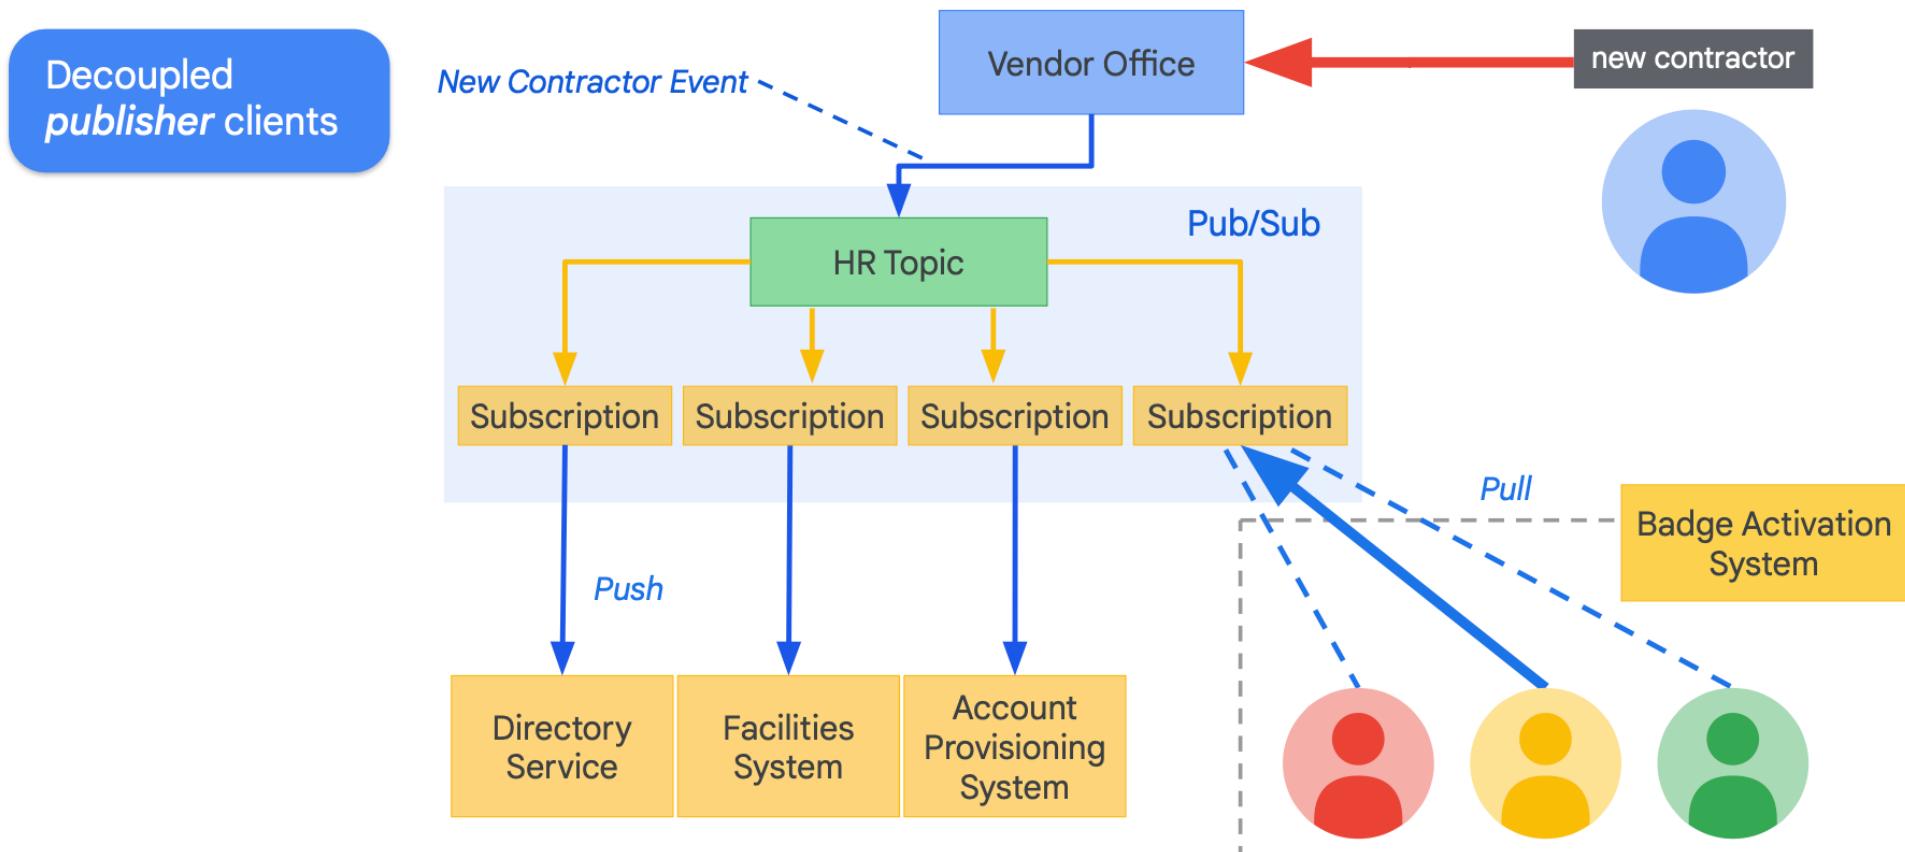
There can be multiple Subscriptions for each Topic



And there can be multiple subscribers per Subscription

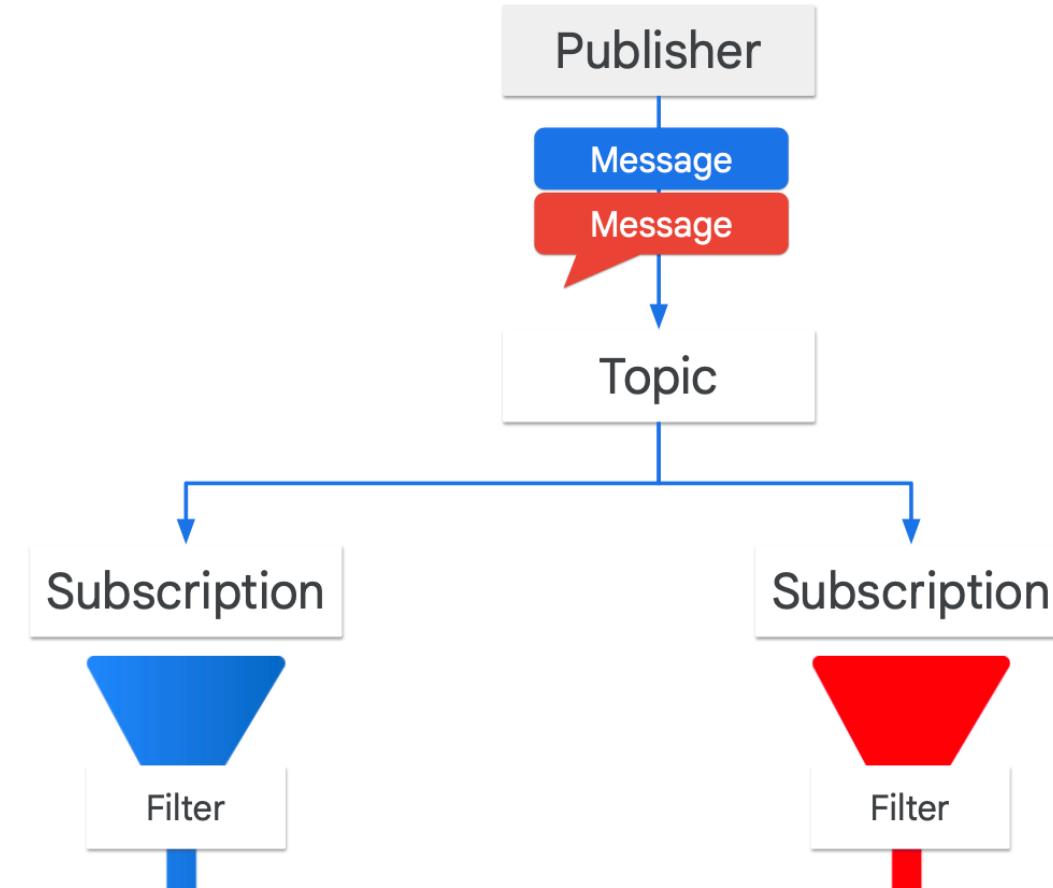


And there can be multiple publishers to the Topic



You can filter messages by their attributes

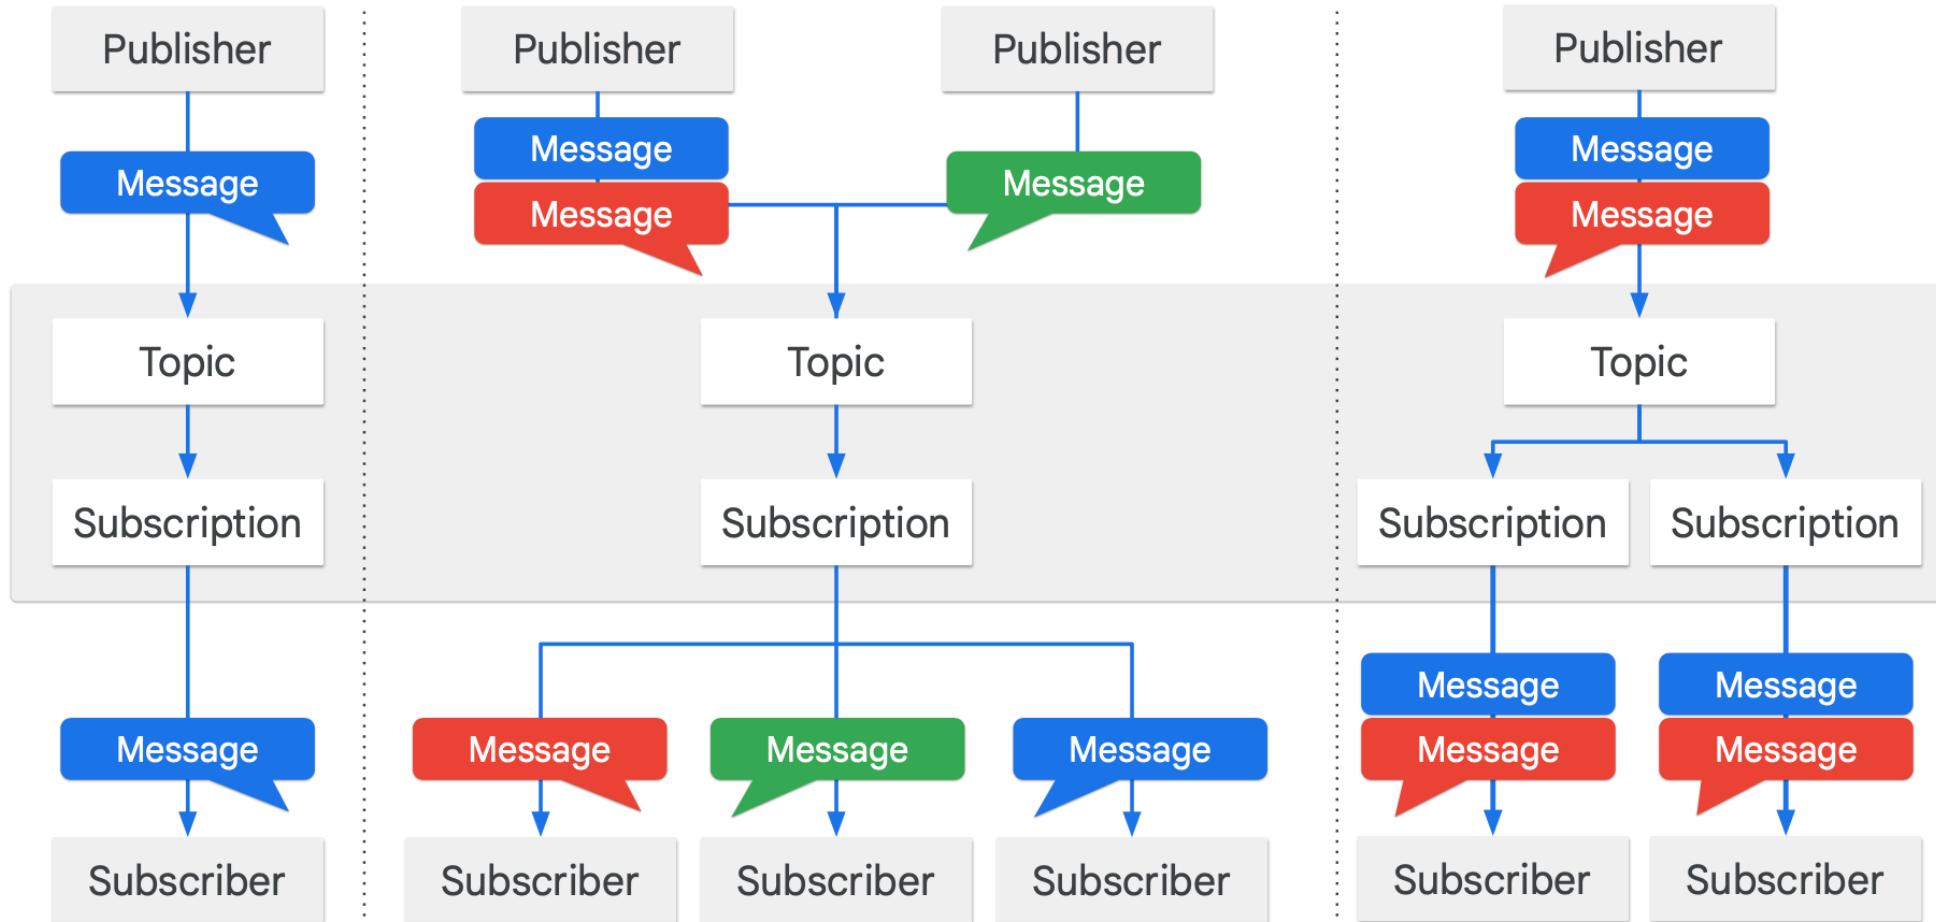
- Filter the Pub/Sub events on the message attributes.
- Configure via the Cloud Console, the gcloud command-line tool, or the Pub/Sub API.



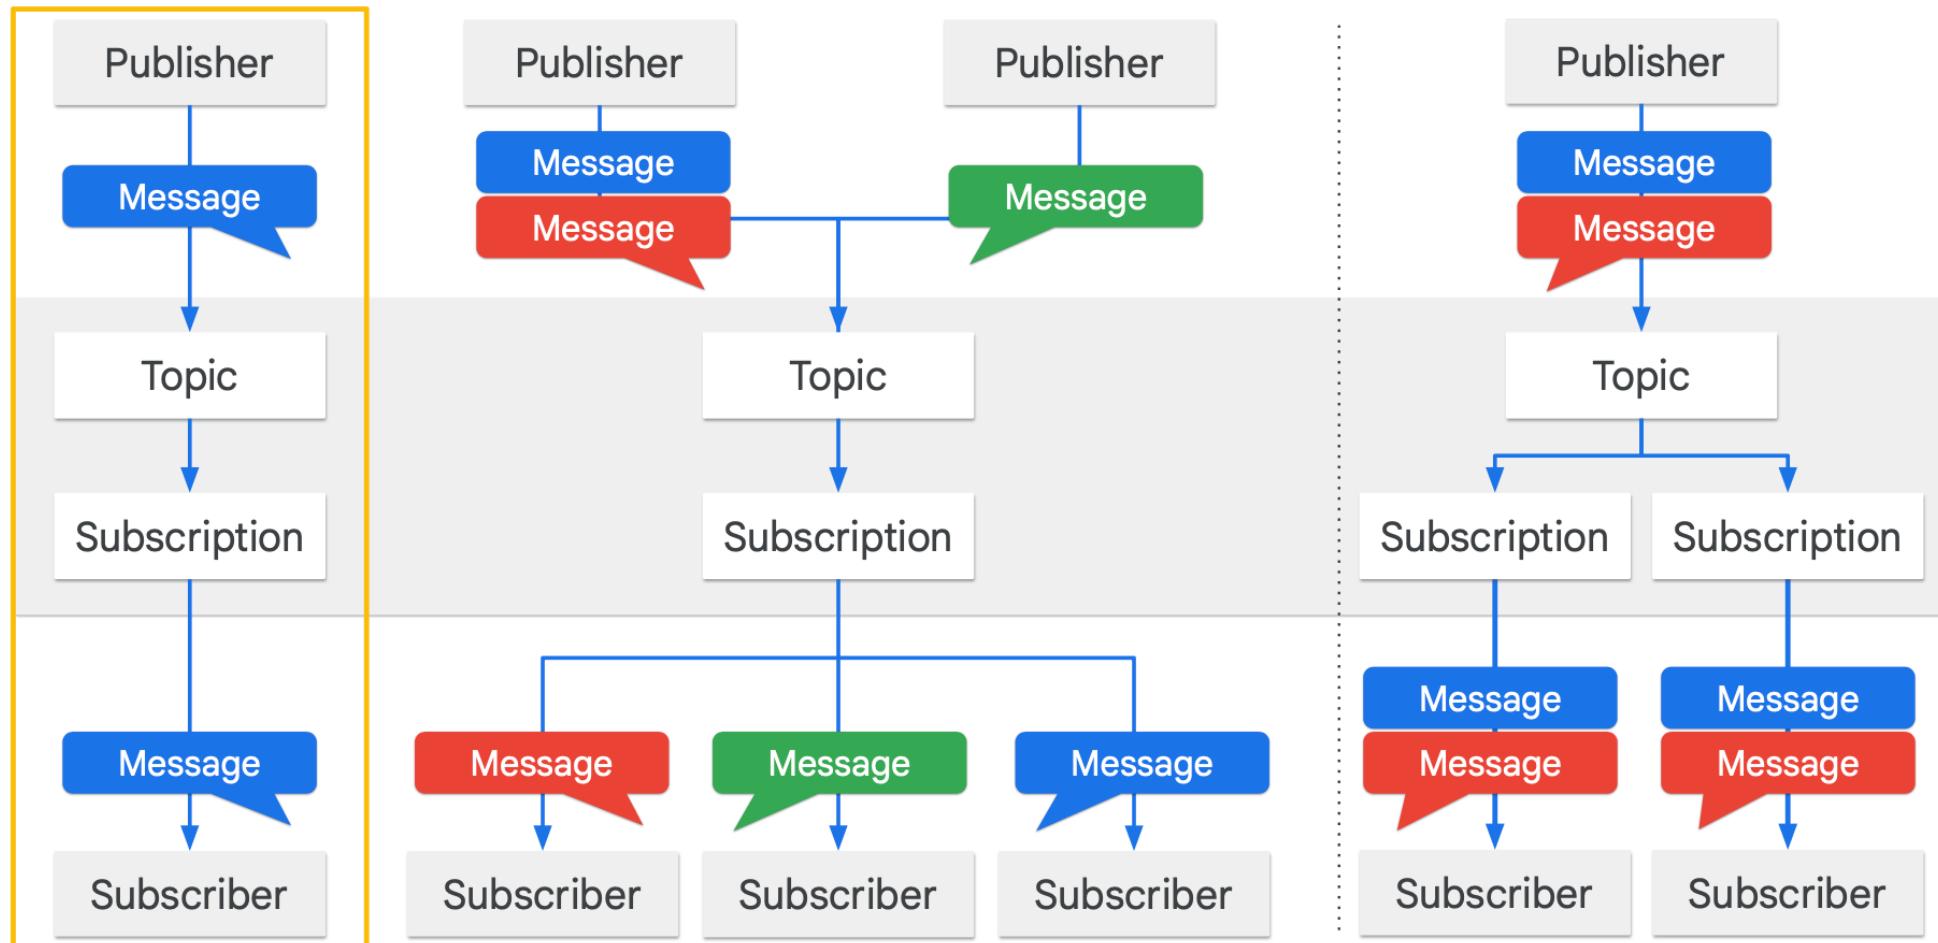


Pub/Sub Push Versus Pull

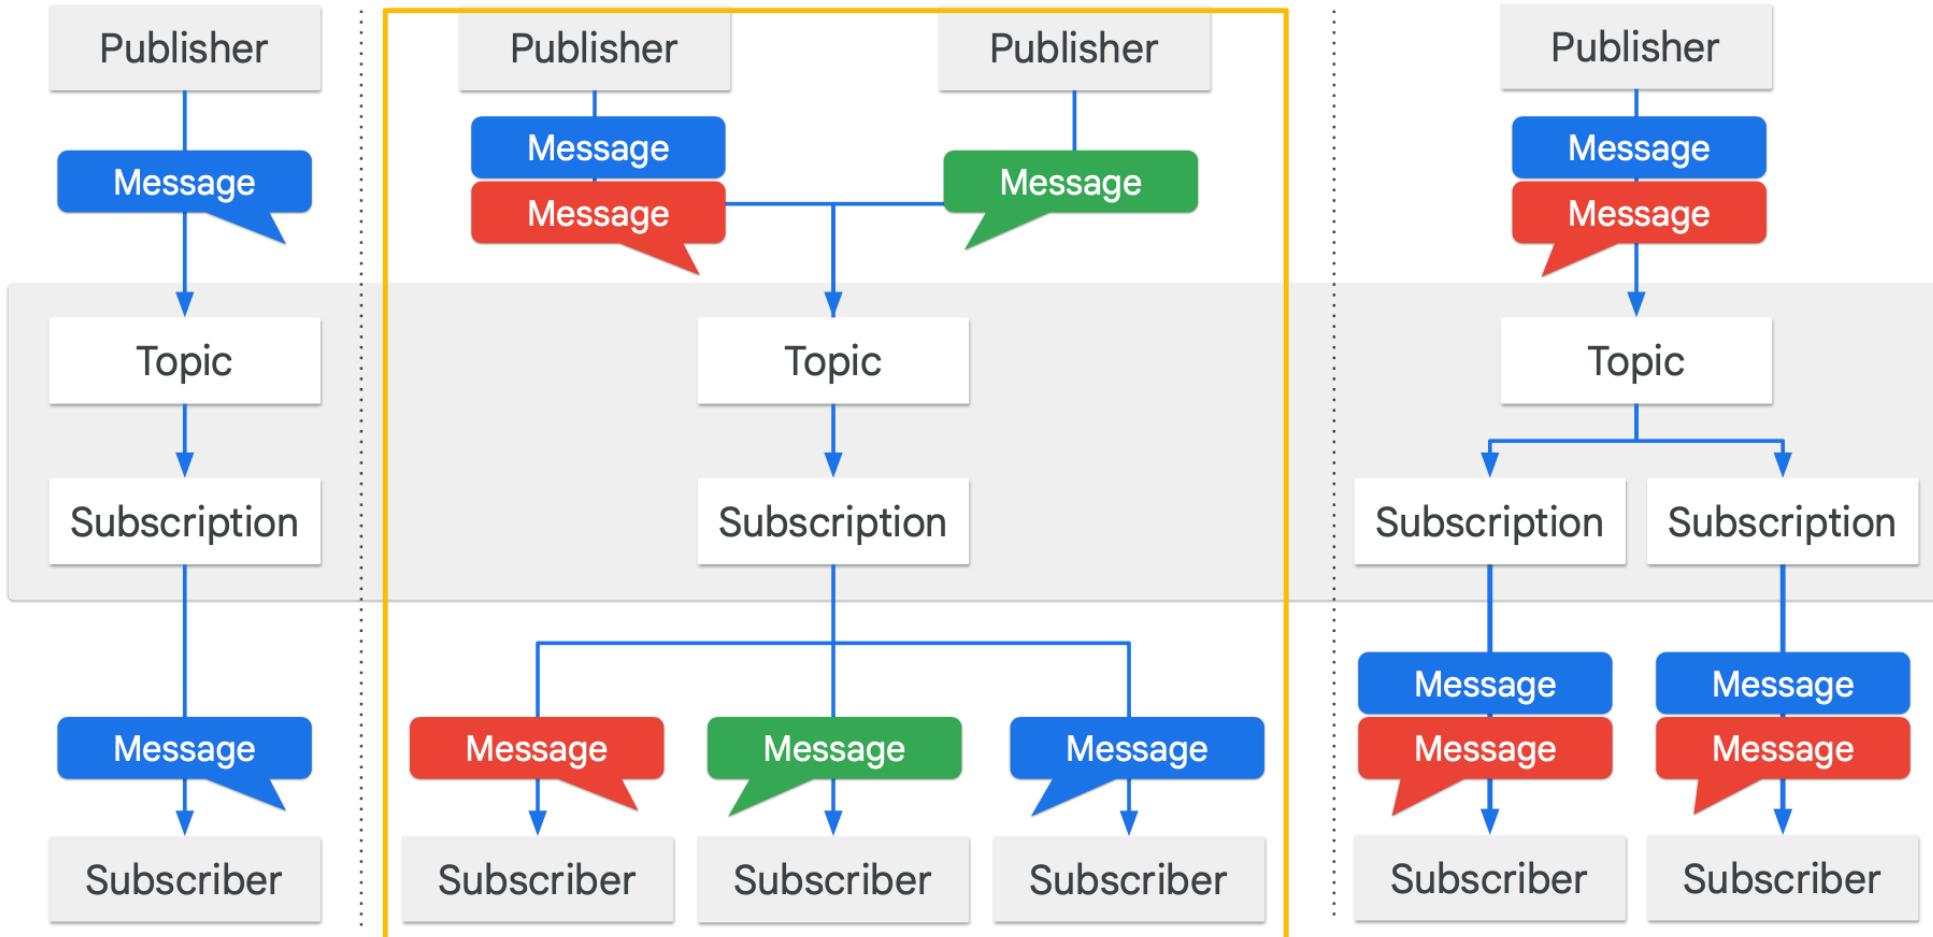
Publish/Subscribe patterns



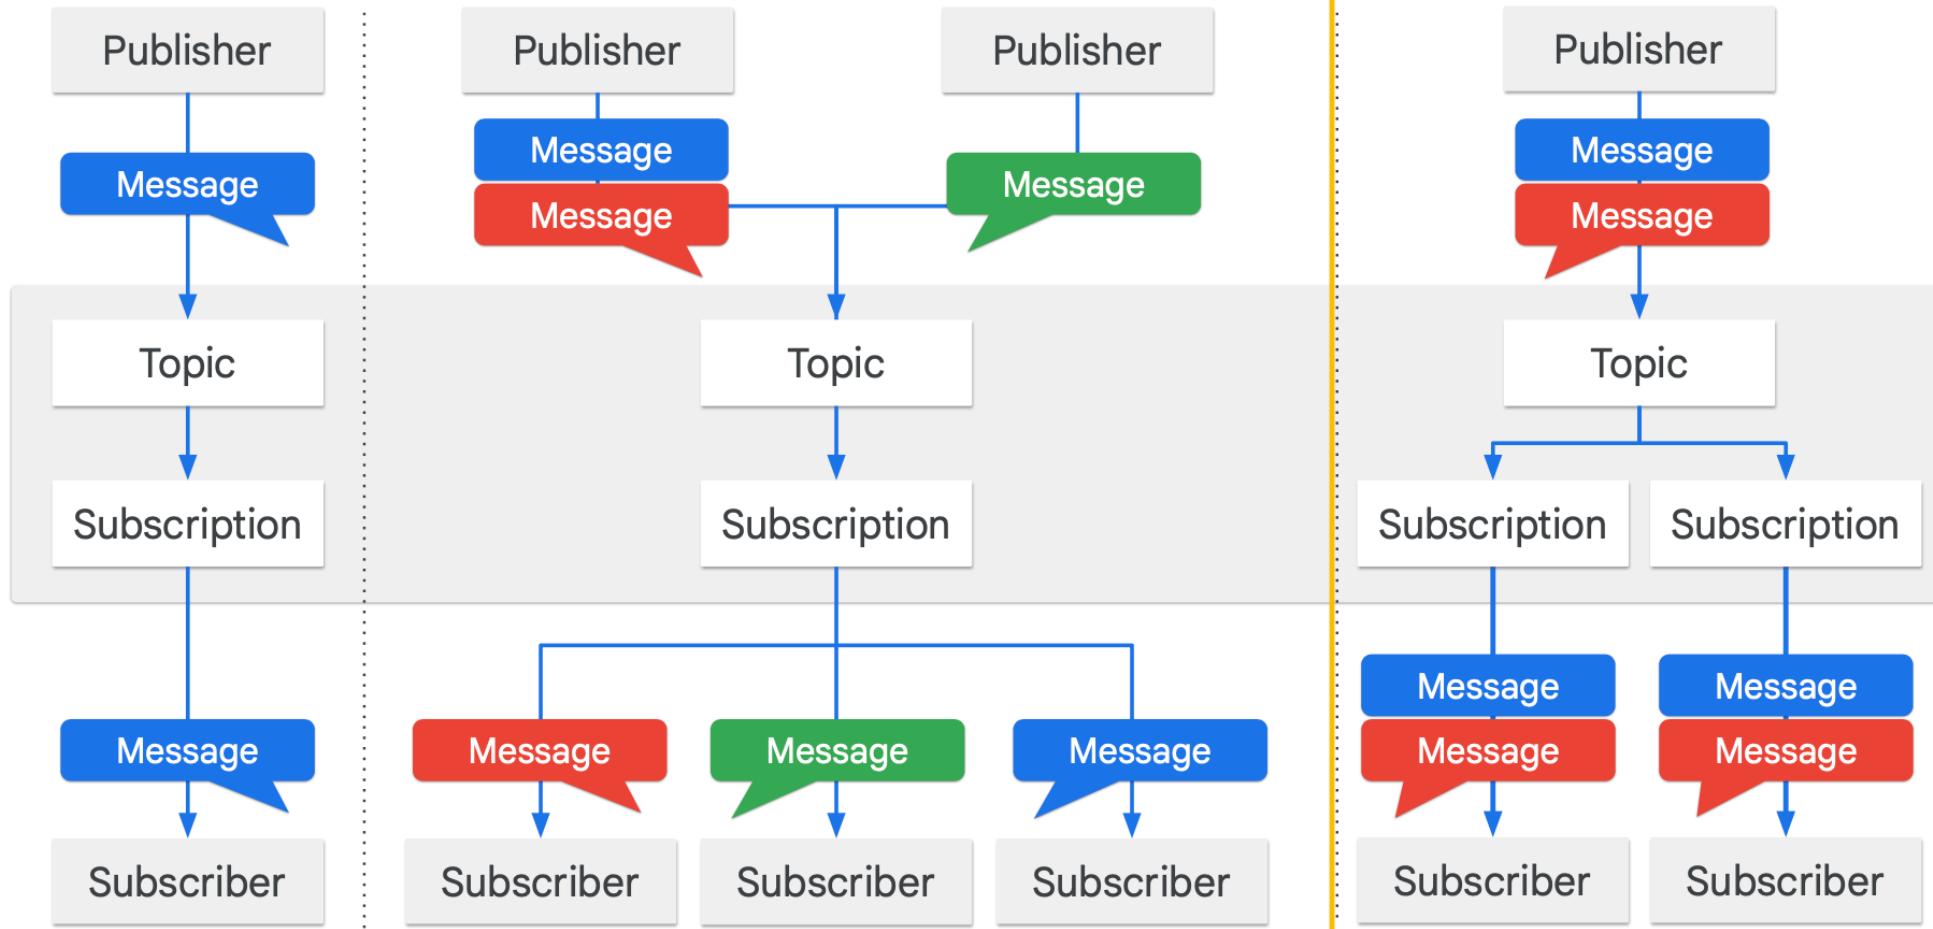
Publish/Subscribe patterns



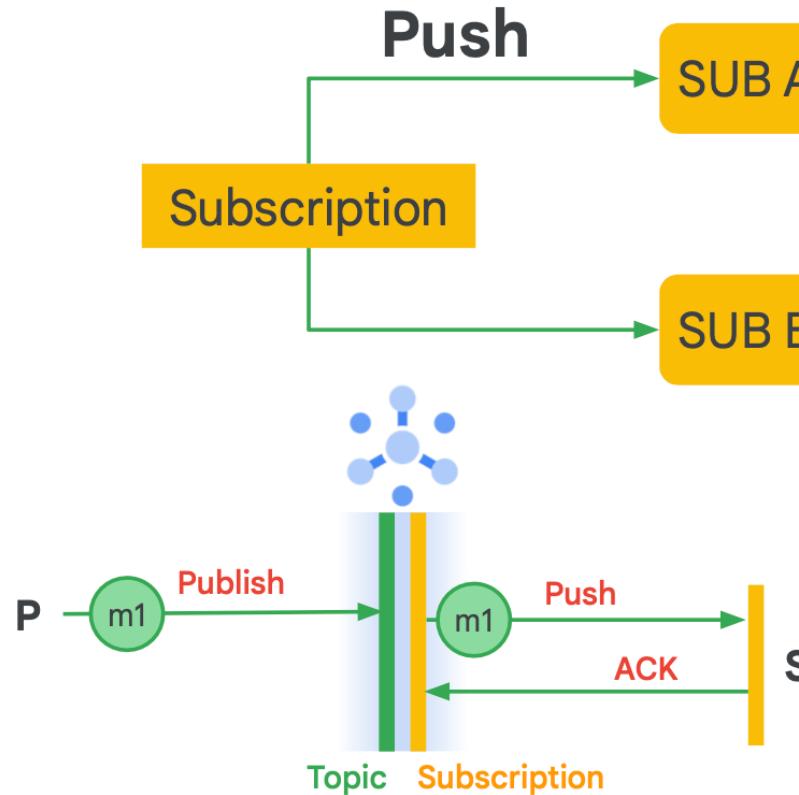
Publish/Subscribe patterns



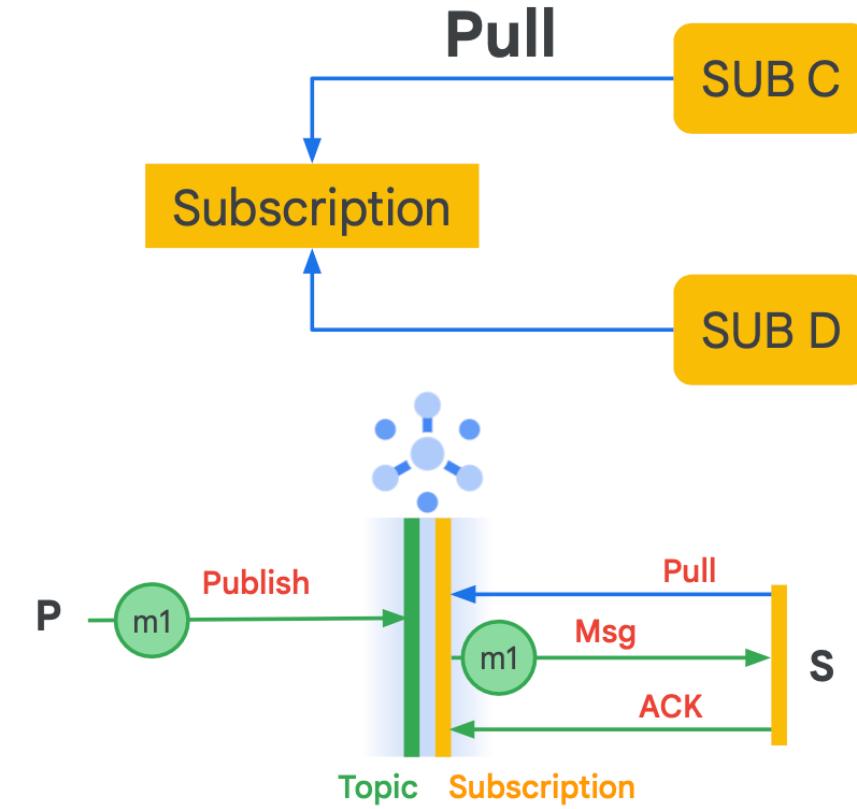
Publish/Subscribe patterns



Pub/Sub provides both Push and Pull delivery



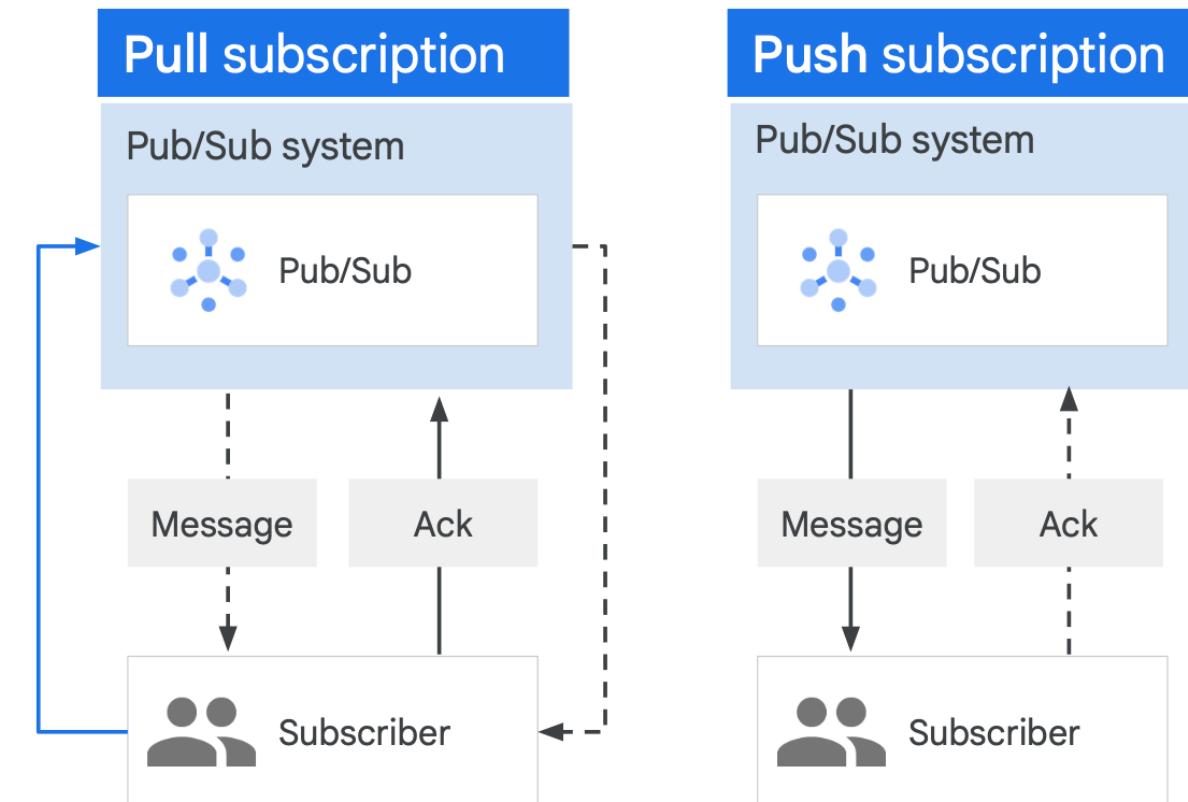
Acknowledgement used for dynamic rate control



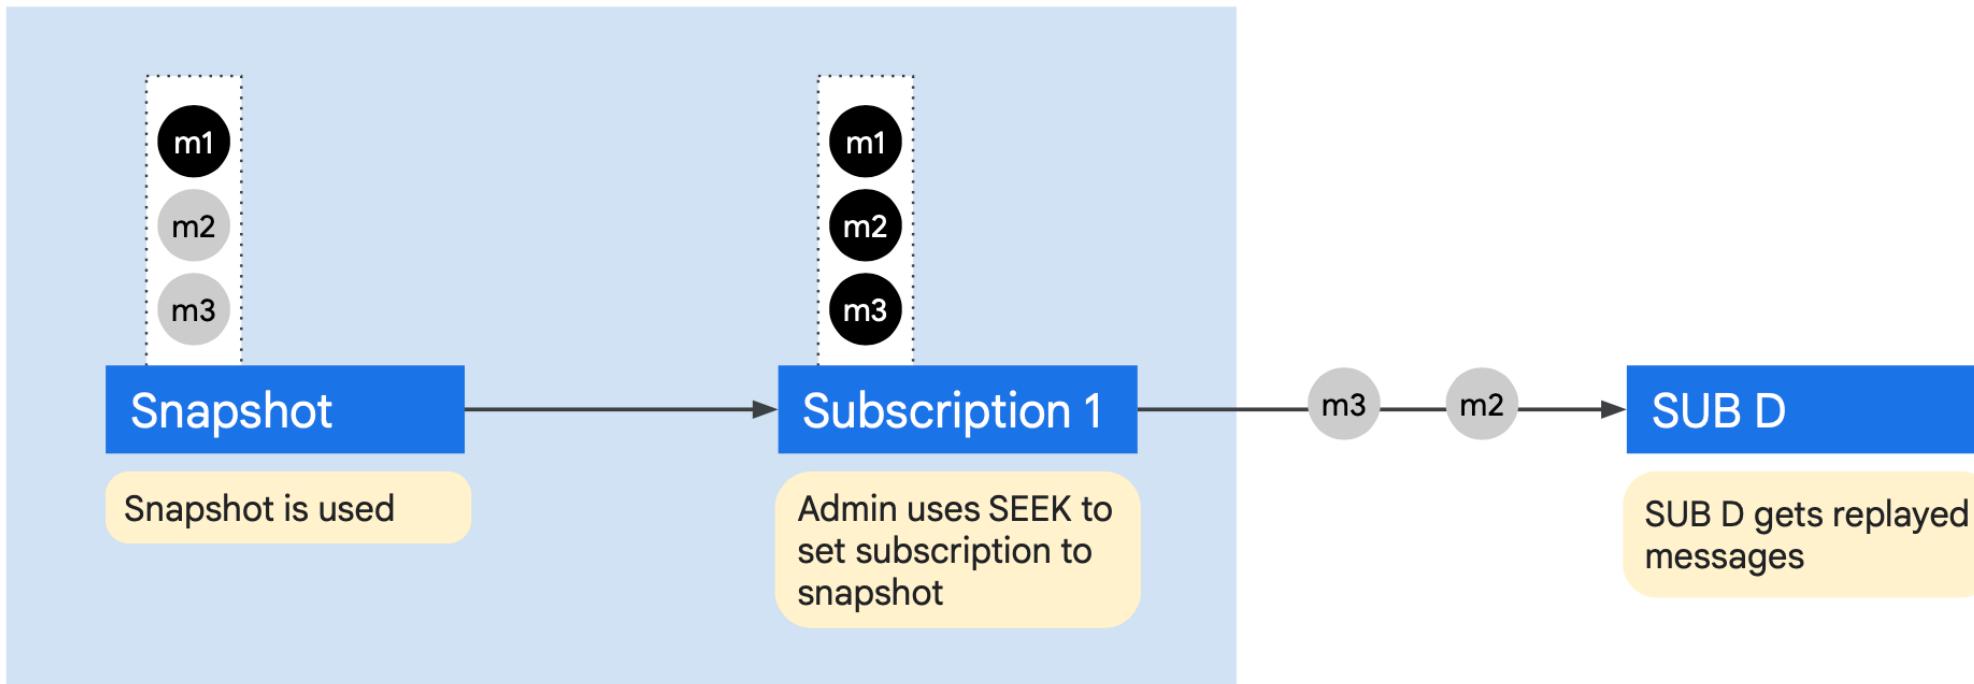
Messages are stored up to 7 days

At least once delivery guarantee

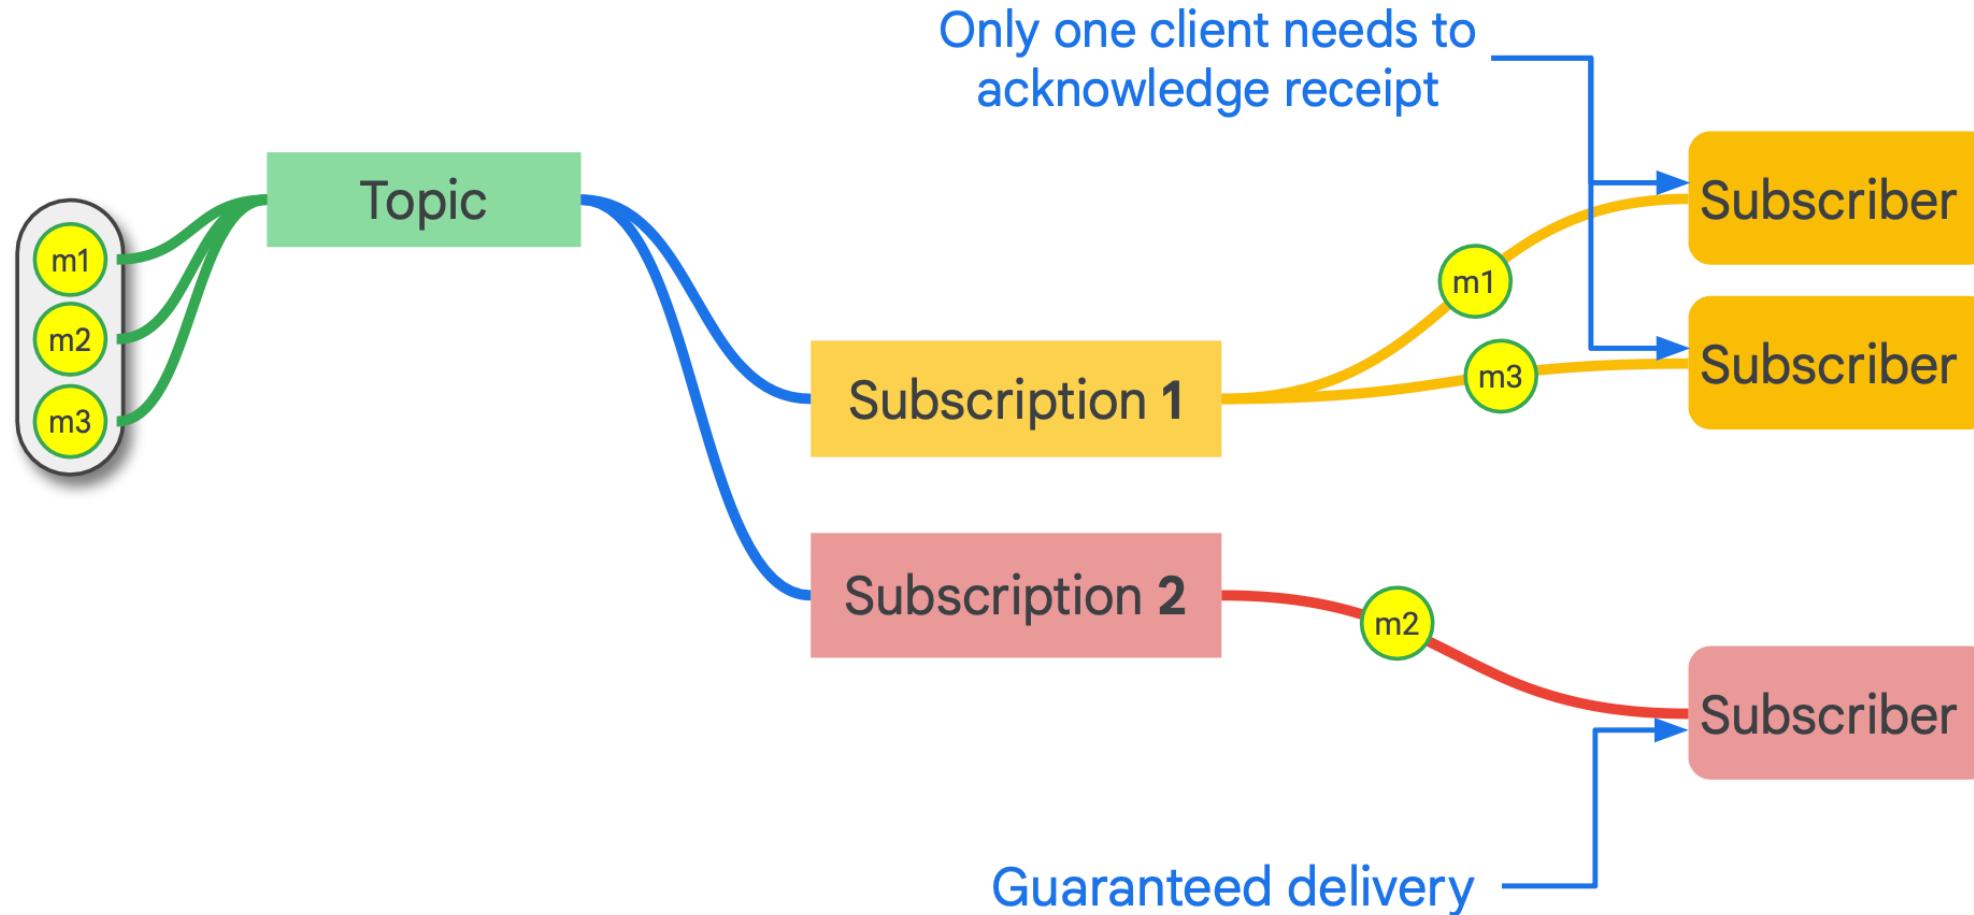
- A subscriber ACKs each message for every subscription
- A message is resent if subscriber takes more than **ackDeadline** to respond
- Messages are stored for up to 7 days
- A subscriber can extend the deadline per message



Message replay



Subscribers can work as a team or separately



03



Publishing with Pub/Sub Code

Publishing with Pub/Sub

```
gcloud pubsub topics create sandiego
```

Create topic

```
gcloud pubsub topics publish sandiego --message "hello"
```

Publish to topic

```
import os
from google.cloud import pubsub_v1
```

```
publisher = pubsub_v1.PublisherClient()
```

```
topic_name = 'projects/{project_id}/topics/{topic}'.format(
    project_id=os.getenv('GOOGLE_CLOUD_PROJECT'),  
          topic='MY_TOPIC_NAME',
)
```

```
publisher.create_topic(topic_name)
publisher.publish(topic_name, b'My first message!', author='dylan')
```

Python

Create a client

Message

Send attribute

Subscribing with Pub/Sub using async pull

```
import os
from google.cloud import pubsub_v1
```

Python

```
subscriber = pubsub_v1.SubscriberClient()
topic_name = 'projects/{project_id}/topics/{topic}'.format(
    project_id=os.getenv('GOOGLE_CLOUD_PROJECT'),
    topic='MY_TOPIC_NAME')
subscription_name = 'projects/{project_id}/subscriptions/{sub}'.format(
    project_id=os.getenv('GOOGLE_CLOUD_PROJECT'),
    sub='MY_SUBSCRIPTION_NAME')
subscriber.create_subscription(
    name=subscription_name, topic=topic_name)
```

Create a client

```
def callback(message):
    print(message.data)
    message.ack()
```

Callback when
message received

```
future = subscriber.subscribe(subscription_name, callback)
```

Pull method
Callback function

Subscribing with Pub/Sub using synchronous pull

```
gcloud pubsub subscriptions create --topic sandiego mySub1
```

Create subscription

```
gcloud pubsub subscriptions pull --auto-ack mySub1
```

Pull subscription

```
import time

from google.cloud import pubsub_v1

subscriber = pubsub_v1.SubscriberClient()

subscription_path = subscriber.subscription_path(project_id, subscription_name)
    'projects/{project_id}/subscriptions/{subscription_name}'<----- subscription_path
format

NUM_MESSAGES = 2
ACK_DEADLINE = 30
SLEEP_TIME = 10

# The subscriber pulls a specific number of messages.
response = subscriber.pull(subscription_path, max_messages=NUM_MESSAGES)
```

Set subscription name

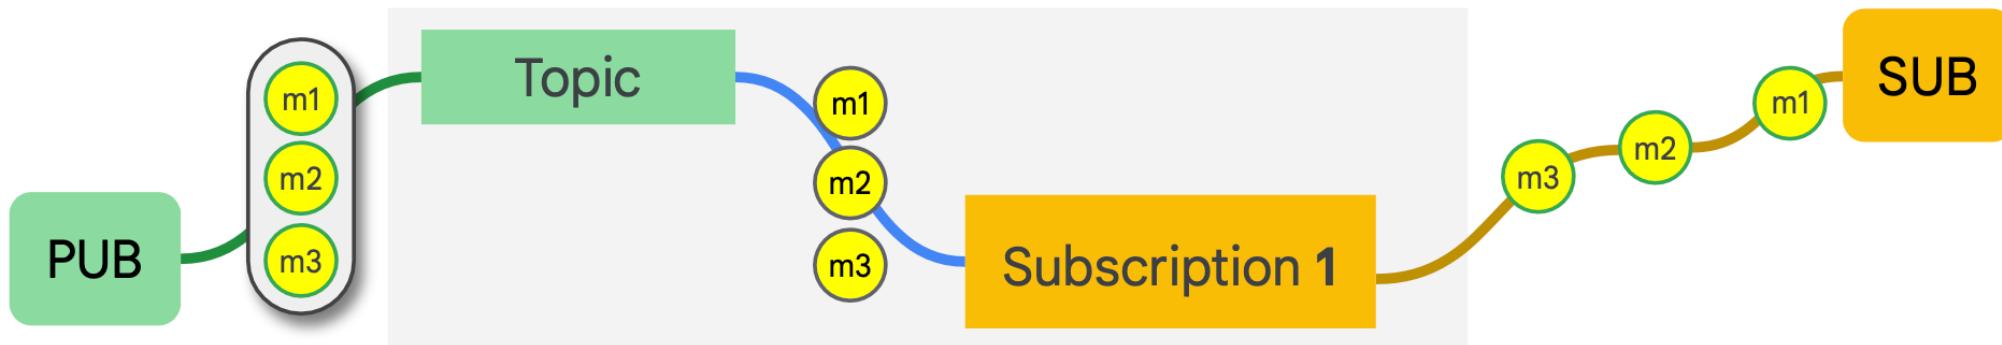
Create a client

subscription_path
format

Subscriber is
non-blocking

Keep the main thread
from exiting to allow it
to process messages
synchronously

By default, the Pub/Sub publishing engine batches messages; turn this off if you desire lower latency



Batching messages: throughput versus latency

Changing the batch settings in Pub/Sub

```
from google.cloud import pubsub  
from google.cloud.pubsub import types  
  
client = pubsub.PublisherClient(  
    batch_settings=BatchSettings(max_messages=500),  
)
```

Python

Change batch setting

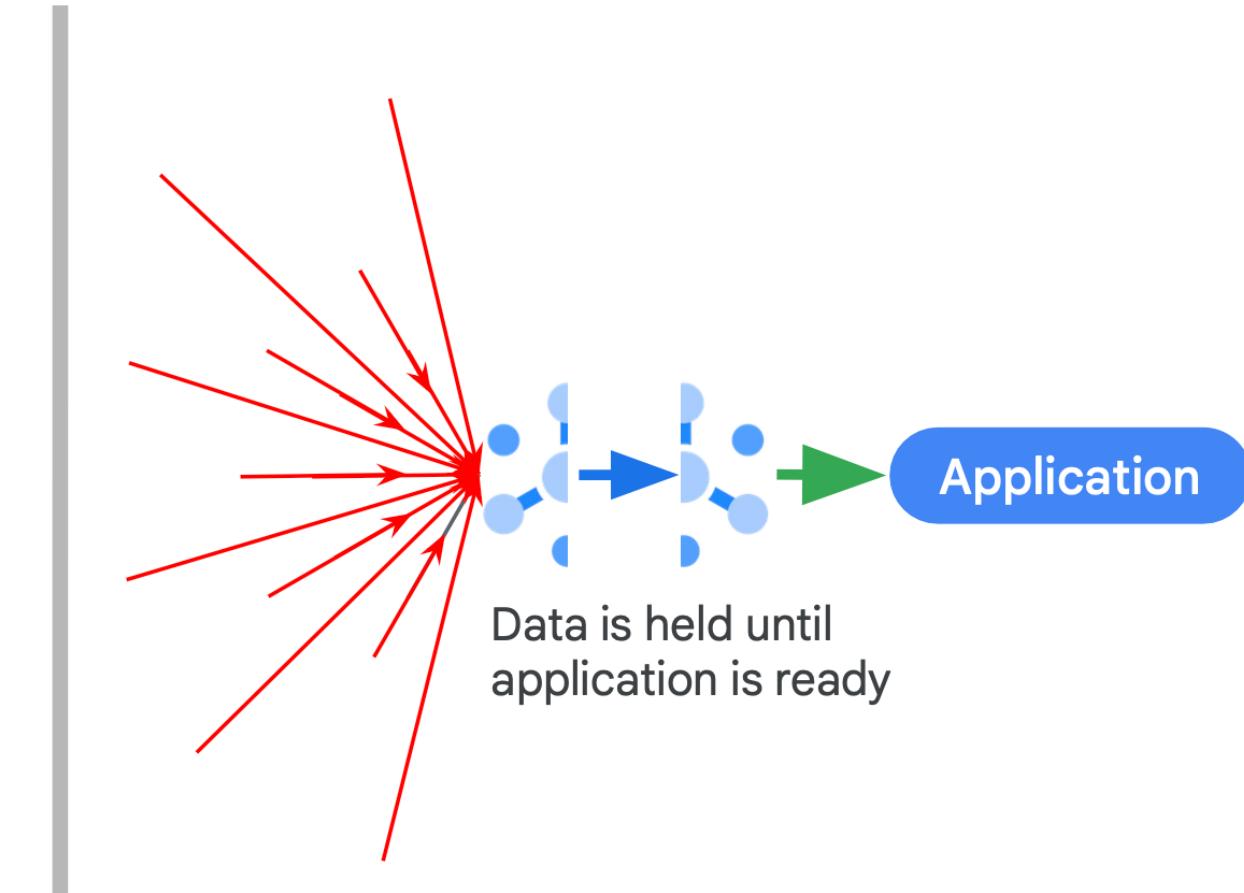
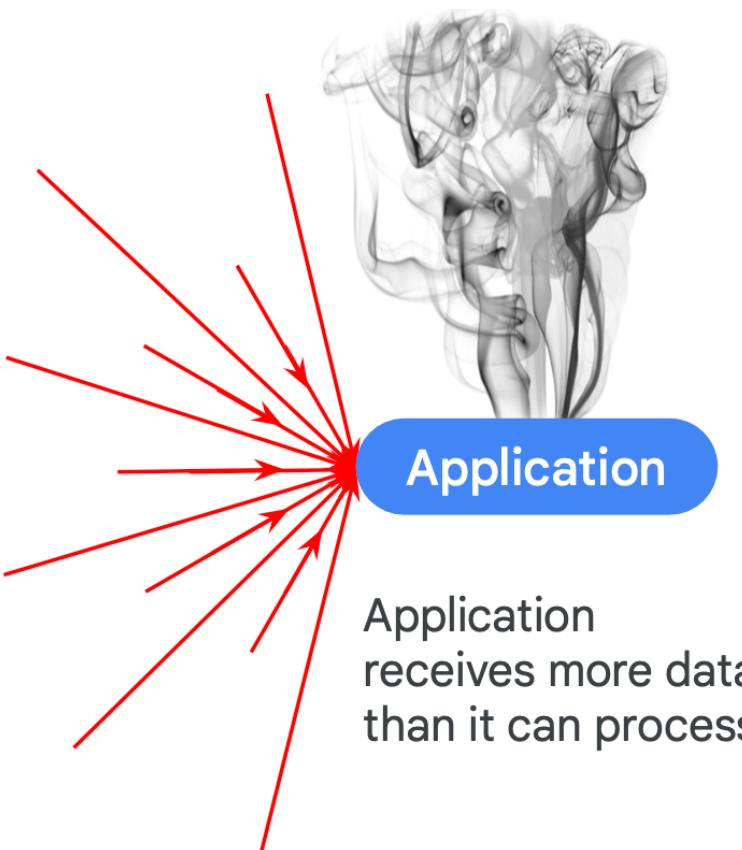
Pub/Sub: Latency, out-of-order, duplication will happen

- If messages have the same ordering key and are in the same region, you can enable message ordering.
- To receive the messages in order, set the message ordering property on the subscription you receive messages from using the Cloud Console, the gcloud command-line tool, or the Pub/Sub API.
- Receiving messages in order might increase latency.

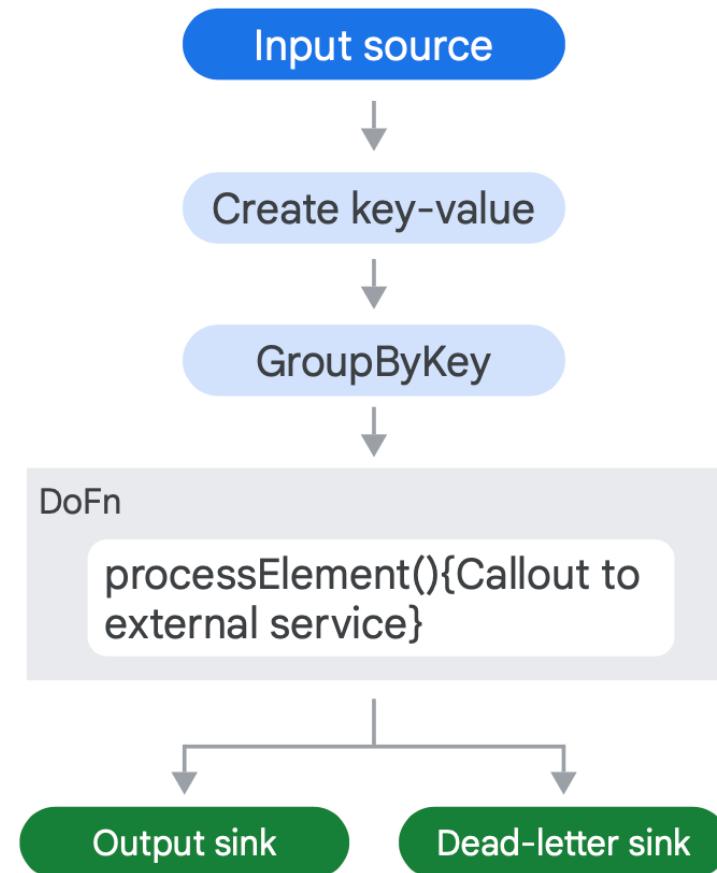
Use Pub/Sub for streaming resilience



Use Pub/Sub for streaming resilience



Dead-letter sinks and error logging



Exponential backoff

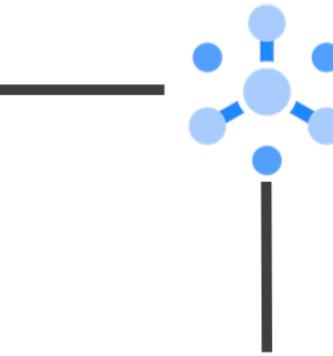
- Pub/Sub lets you configure an exponential backoff policy for better flow control.
- The idea behind exponential backoff is to add progressively longer delays between retry attempts.
- To create a new subscription with an exponential backoff retry policy, run the `gcloud pubsub create` command or use the Cloud console.

Security, monitoring, and logging for Pub/Sub

Cloud Logging
metrics

pubsub_topic

pubsub_subscription



Cloud Audit Logs

Admin = on

Data Access = off

Authentication

Service accounts
User accounts



Access control

Pub/Sub roles per topic

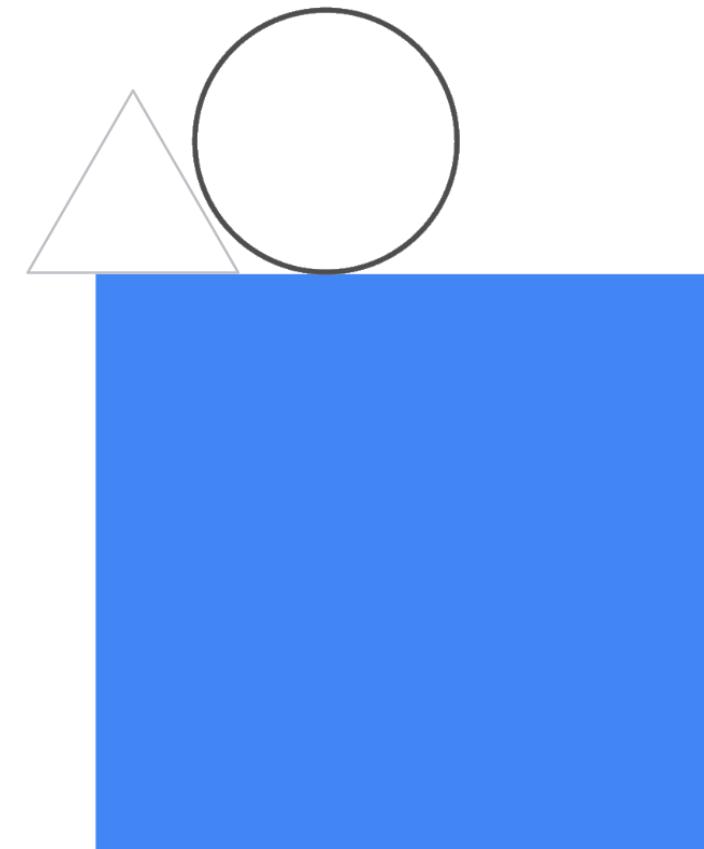
Summary

Latency, out-of-order, duplication will happen

Pub/Sub with Dataflow: Exactly once, ordered processing

Lab Intro

Streaming Data Processing:
Publish Streaming Data
into Pub/Sub



Lab objectives

- 01 Create a Pub/Sub topic and subscription
- 02 Simulate your traffic sensor data into Pub/Sub

